

# InternImage: Exploring Large-Scale Vision Foundation Models with Deformable Convolutions

DEFORMABLE CONVOLUTIONS







발표자: 김병현

이미지처리팀:  
강인하, 김현진, 안종식, 이주영, 이희재, 현청천

# InterImage

## ❖ State-of-the-Art (COCO test-dev) Backbone Network

📦 Released by OpenGVLab

Rank	Model	box AP	AP50	AP75	APs	APM	APL	Params (M)	Extra Training Data	Paper	Code	Result	Year	Tags
1	InternImage-DCNv3-H (M3I Pre-training)	65.4							×	InternImage: Exploring Large- Scale Vision Foundation Models with Deformable Convolutions			2022	<div>DCN</div> <div>Deformable Convolution</div> <div>DINO</div> <div>Giant</div>
2	M3I Pre-training (InternImage-H)	65.4							×	Towards All-in- one Pre-training via Maximizing Multi-modal Mutual Information			2022	
3	EVA	64.7	81.9	71.7	48.5	67.7	77.9		×	EVA: Exploring the Limits of Masked Visual Representation Learning at Scale			2022	



# Introduction

---

- ❖ Success of Transformers in Computer Vision Tasks
- ❖ CNN-based foundation models can also achieve comparable or even better performance than ViTs when equipped with similar operator-/architecture-level designs, scaling-up parameters, and massive data.



# Introduction

---

## ❖ Gap between CNNs and ViTs

### ▣ Operator Level

- Long-range dependency
- Adaptive spatial aggregation

### ▣ Architecture View

- Advanced components
  - Layer Normalization
  - Feed Forward Networks
  - GELU


### ▣ Recent Long-Range CNNs

- Very large kernels (31x31)
- Gap with SOTA ViTs

# Introduction


## ❖ Comparison of different core operators

★ query pixels    ■ response pixels with **fixed** weights  
■ ■ ■ response pixels with **adaptive** weights




(a) global attention

✓ long-range dependence  
✓ adaptive spatial aggregation  
✗ computation/memory efficient




(b) local attention

✗ long-range dependence  
✓ adaptive spatial aggregation  
✓ computation/memory efficient



(c) large kernel

✓ long-range dependence  
✗ adaptive spatial aggregation  
✓ computation/memory efficient

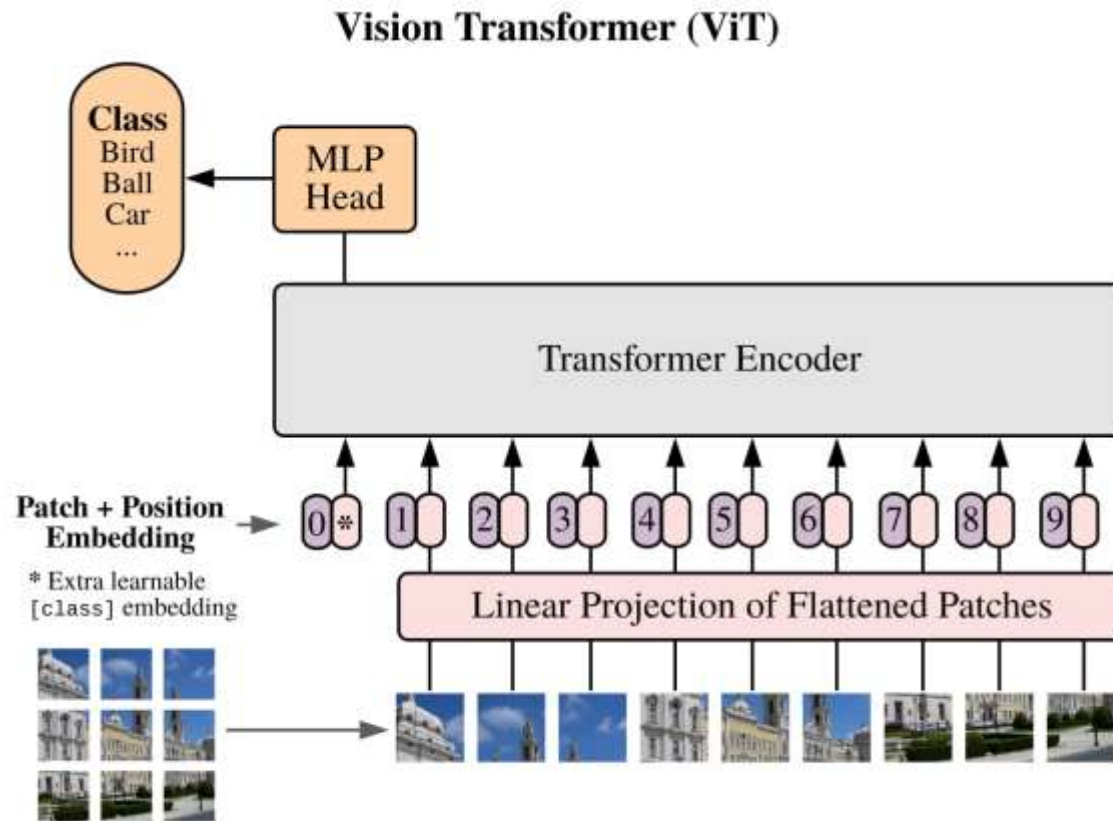


(d) dynamic sparse kernel (ours)

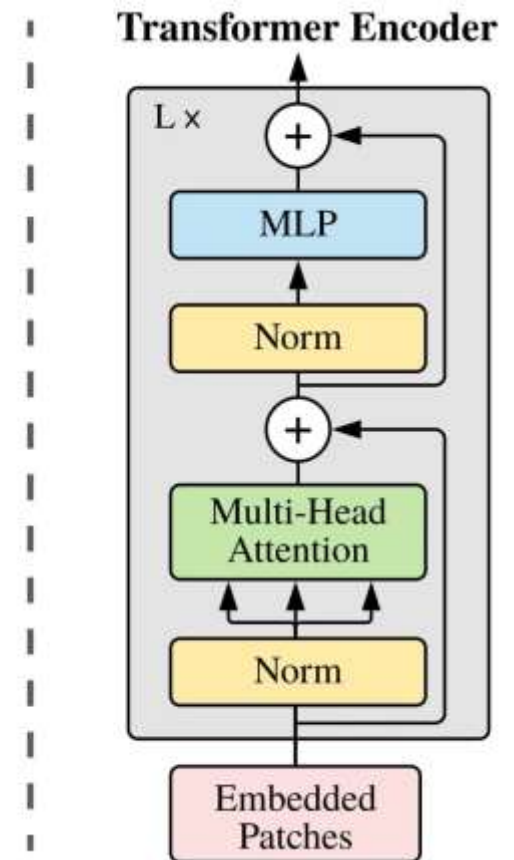
✓ long-range dependence  
✓ adaptive spatial aggregation  
✓ computation/memory efficient

# Introduction

## ❖ Global Attention: Vision Transformer

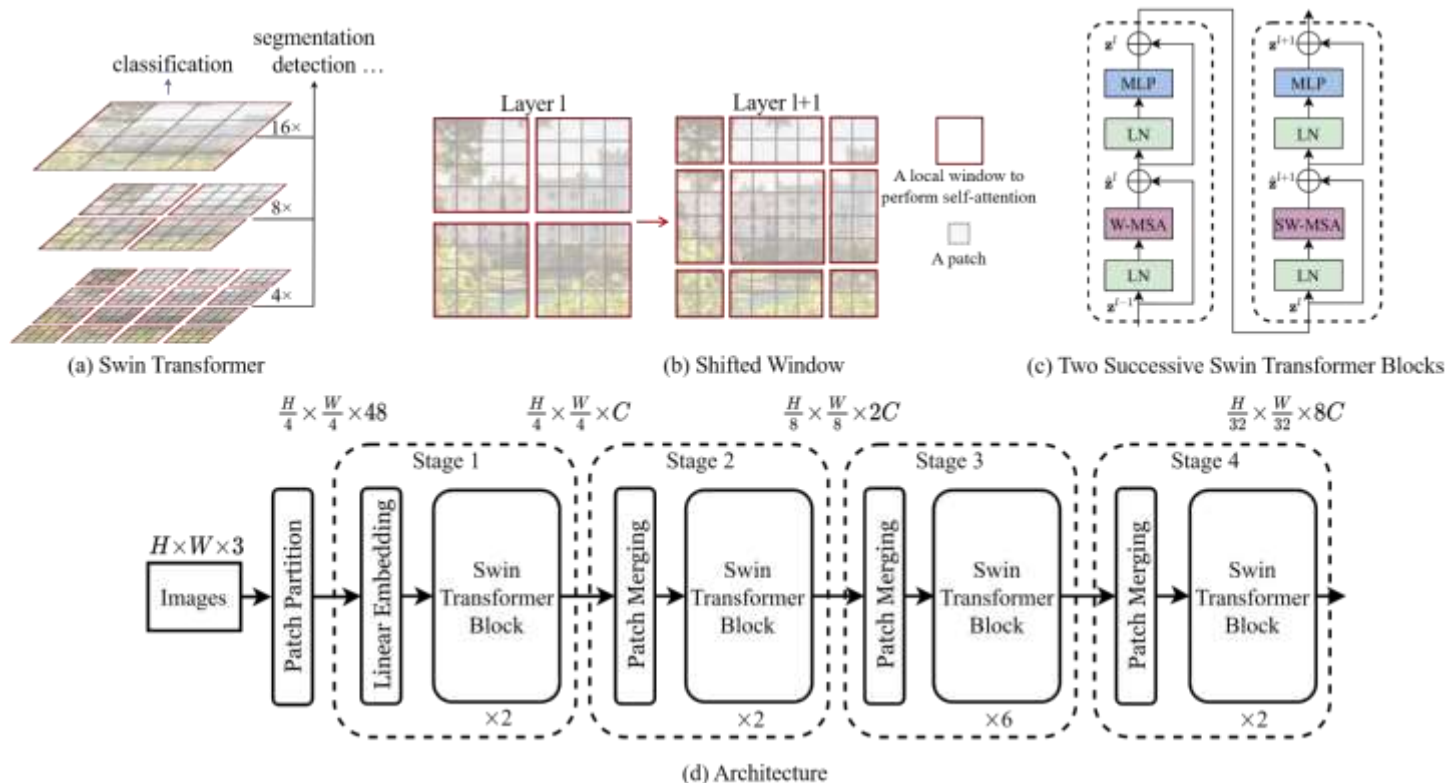


Architecture of ViT



# Introduction

## ❖ Local Attention: Swin Transformer



Architecture of Swin Transformer

# Introduction

## ❖ Large Kernel: SLaK

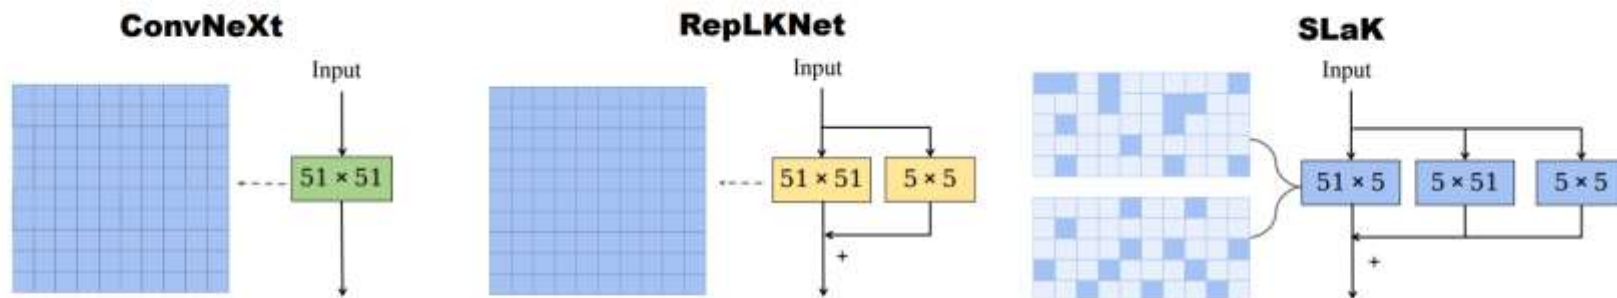


Figure 1: Large depth-wise kernel (e.g.,  $51 \times 51$ ) paradigms of ConvNeXt, RepLKNet, and SLaK. Dark blue squares refer to the dense weights in convolutional kernels. Light blue squares refer to the sparse weights in convolutional kernels.





# Introduction

---

## ❖ Concentration on CNN-based Model

### ▣ *InterImage*

- Brand-New CNN-based Backbone Network
- Characteristics
  - Dynamic sparse convolutional layer
    - » Only with 3x3 kernels
    - » Adaptive spatial aggregation
    - » Reduce inductive bias
    - » Low computational cost compared to large convolutional layers
  - Overall Architecture of ViT

# Introduction

## ❖ Contributions

- ❑ 1<sup>st</sup> CNN-based backbone with more than 1 billion params.
- ❑ Add long-range dependencies and adaptive spatial aggregation with 3x3 DCN
- ❑ SOTA accuracy in COCO dataset

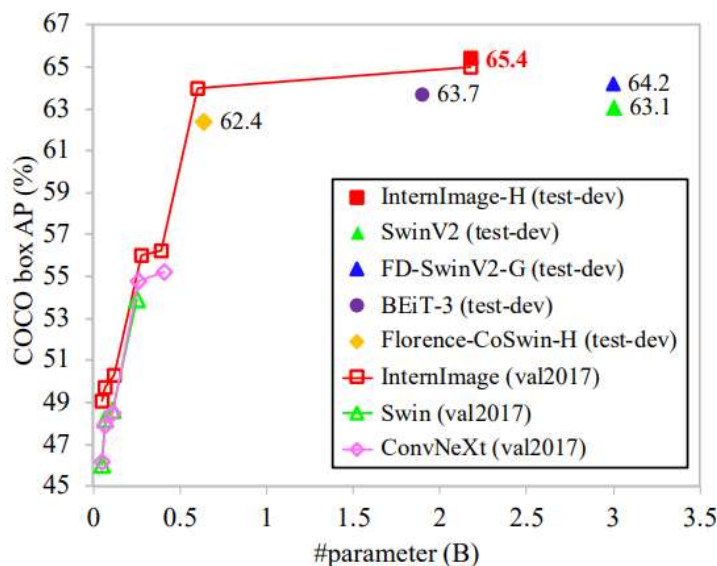


Figure 2. **Performance comparison on COCO of different backbones.** The proposed InternImage-H achieves a new record 65.4 box AP on COCO test-dev, significantly outperforming state-of-the-art CNNs and large-scale ViTs.



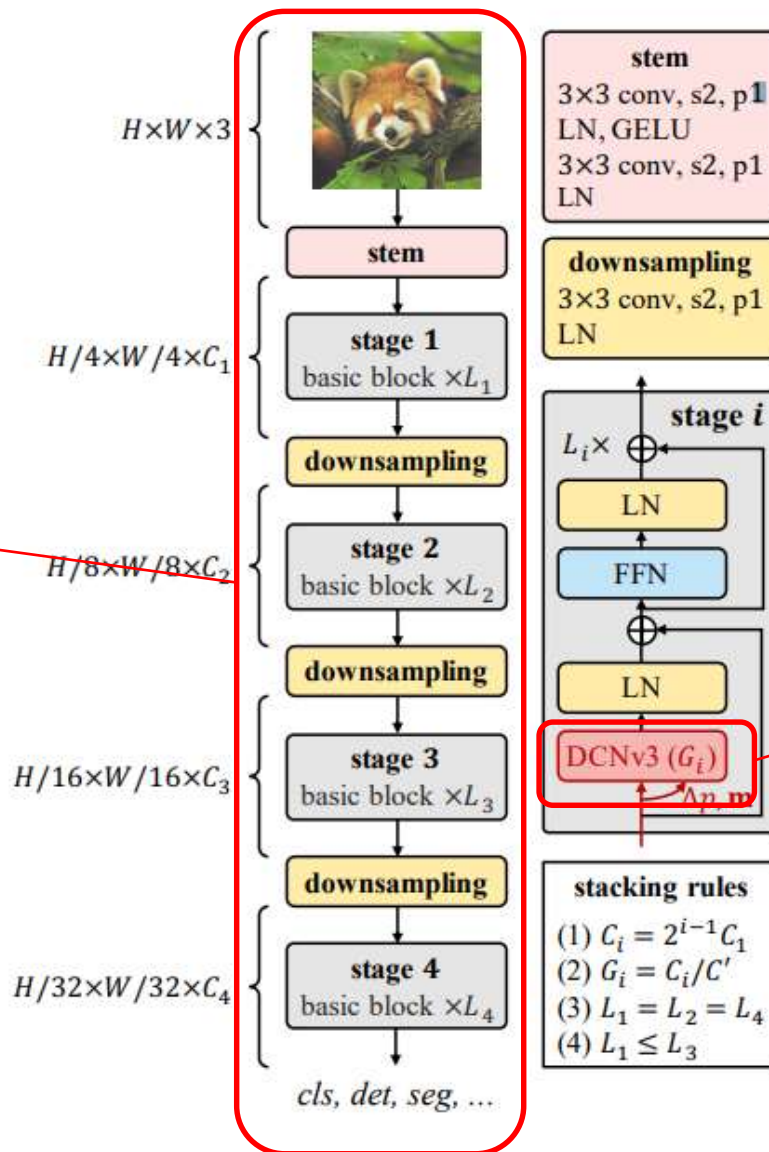
## Q & A

---

Q & A

# Overall Architecture

## 2. Architecture Design



1. Deformable Convolutional Layer V3



# Deformable Convolutional Layer v3

---

## ❖ Revisiting DCNv2


**Revisiting DCNv2.** A straightforward way to bridge the gap between convolution and MHSA is to introduce long-range dependencies and adaptive spatial aggregation into regular convolutions. Let us start with DCNv2 [28], which is a general variant of regular convolution. Given an input  $\mathbf{x} \in \mathbb{R}^{C \times H \times W}$  and current pixel  $p_0$ , DCNv2 can be formulated as:

$$\mathbf{y}(p_0) = \sum_{k=1}^K \mathbf{w}_k \mathbf{m}_k \mathbf{x}(p_0 + p_k + \Delta p_k), \quad (1)$$

where  $K$  represents the total number of sampling points, and  $k$  enumerates the sampling point.  $\mathbf{w}_k \in \mathbb{R}^{C \times C}$  denotes the projection weights of the  $k$ -th sampling point, and  $\mathbf{m}_k \in \mathbb{R}$  represents the modulation scalar of the  $k$ -th sampling point, which is normalized by sigmoid function.  $p_k$  denotes the  $k$ -th location of the pre-defined grid sampling  $\{(-1, -1), (-1, 0), \dots, (0, +1), \dots, (+1, +1)\}$  as in regular convolutions, and  $\Delta p_k$  is the offset corresponding to the  $k$ -th grid sampling location. We see from the equation that (1) for long-range dependencies, the sampling offset  $\Delta p_k$  is flexible and able to interact with short- or

# Deformable Convolutional Layer v3

## Quick survey of Deformable ConvNets



Regular convolution

$$y(p_0) = \sum_{p_n \in \mathcal{R}} w(p_n) \cdot x(p_0 + p_n)$$

Deformable convolution

$$y(p_0) = \sum_{p_n \in \mathcal{R}} w(p_n) \cdot x(p_0 + p_n + \Delta p_n)$$

where  $\Delta p_n$  is generated by a sibling branch of regular convolution



The grid sampling locations of standard convolution are each offset by displacements learned with respect to the preceding feature maps.



# Deformable Convolutional Layer v3

## Modulated Deformable Modules: DCN-v2

$$y(p) = \sum_{k=1}^K w_k \cdot x(p + p_k + \Delta p_k) \cdot \Delta m_k$$



With modulation, the Deformable ConvNets modules can not only adjust offsets in perceiving input features, but also modulate the input feature **amplitudes/weights** from different spatial locations.



# Deformable Convolutional Layer v3

---

1. Sharing weights among convolutional neurons.
  - Heavy computational cost of DCNv2
    - independent linear projection weights
    - memory complexity is linear with the total number of sampling points
  - To remedy this problem, we borrow the idea from the separable convolution and detach the original convolution weights into depth-wise and point-wise parts
2. Introducing multi-group mechanism
  - Split the spatial aggregation process into  $G$  groups
3. Normalizing modulation scalars along sampling points
  - Change element-wise sigmoid normalization to softmax normalization along sample points.





# Deformable Convolutional Layer v3

---


Combining the aforementioned modifications, the extended DCNv2, marked as DCNv3, can be formulated as Eqn. (2).

$$\mathbf{y}(p_0) = \sum_{g=1}^G \sum_{k=1}^K \mathbf{w}_g \mathbf{m}_{gk} \mathbf{x}_g(p_0 + p_k + \Delta p_{gk}), \quad (2)$$

where  $G$  denotes the total number of aggregation groups. For the  $g$ -th group,  $\mathbf{w}_g \in \mathbb{R}^{C \times C'}$ ,  $\mathbf{m}_{gk} \in \mathbb{R}$  denote the location-irrelevant projection weights of the group, where  $C' = C/G$  represents the group dimension.  $\mathbf{m}_{gk} \in \mathbb{R}$  denotes the modulation scalar of the  $k$ -th sampling point in the  $g$ -th group, normalized by the softmax function along the dimension  $K$ .  $\mathbf{x}_g \in \mathbb{R}^{C' \times H \times W}$  represents the sliced input feature map.  $\Delta p_{gk}$  is the offset corresponding to the grid sampling location  $p_k$  in the  $g$ -th group.


# Deformable Convolutional Layer v3

## ❖ Normal Convolutional Layer



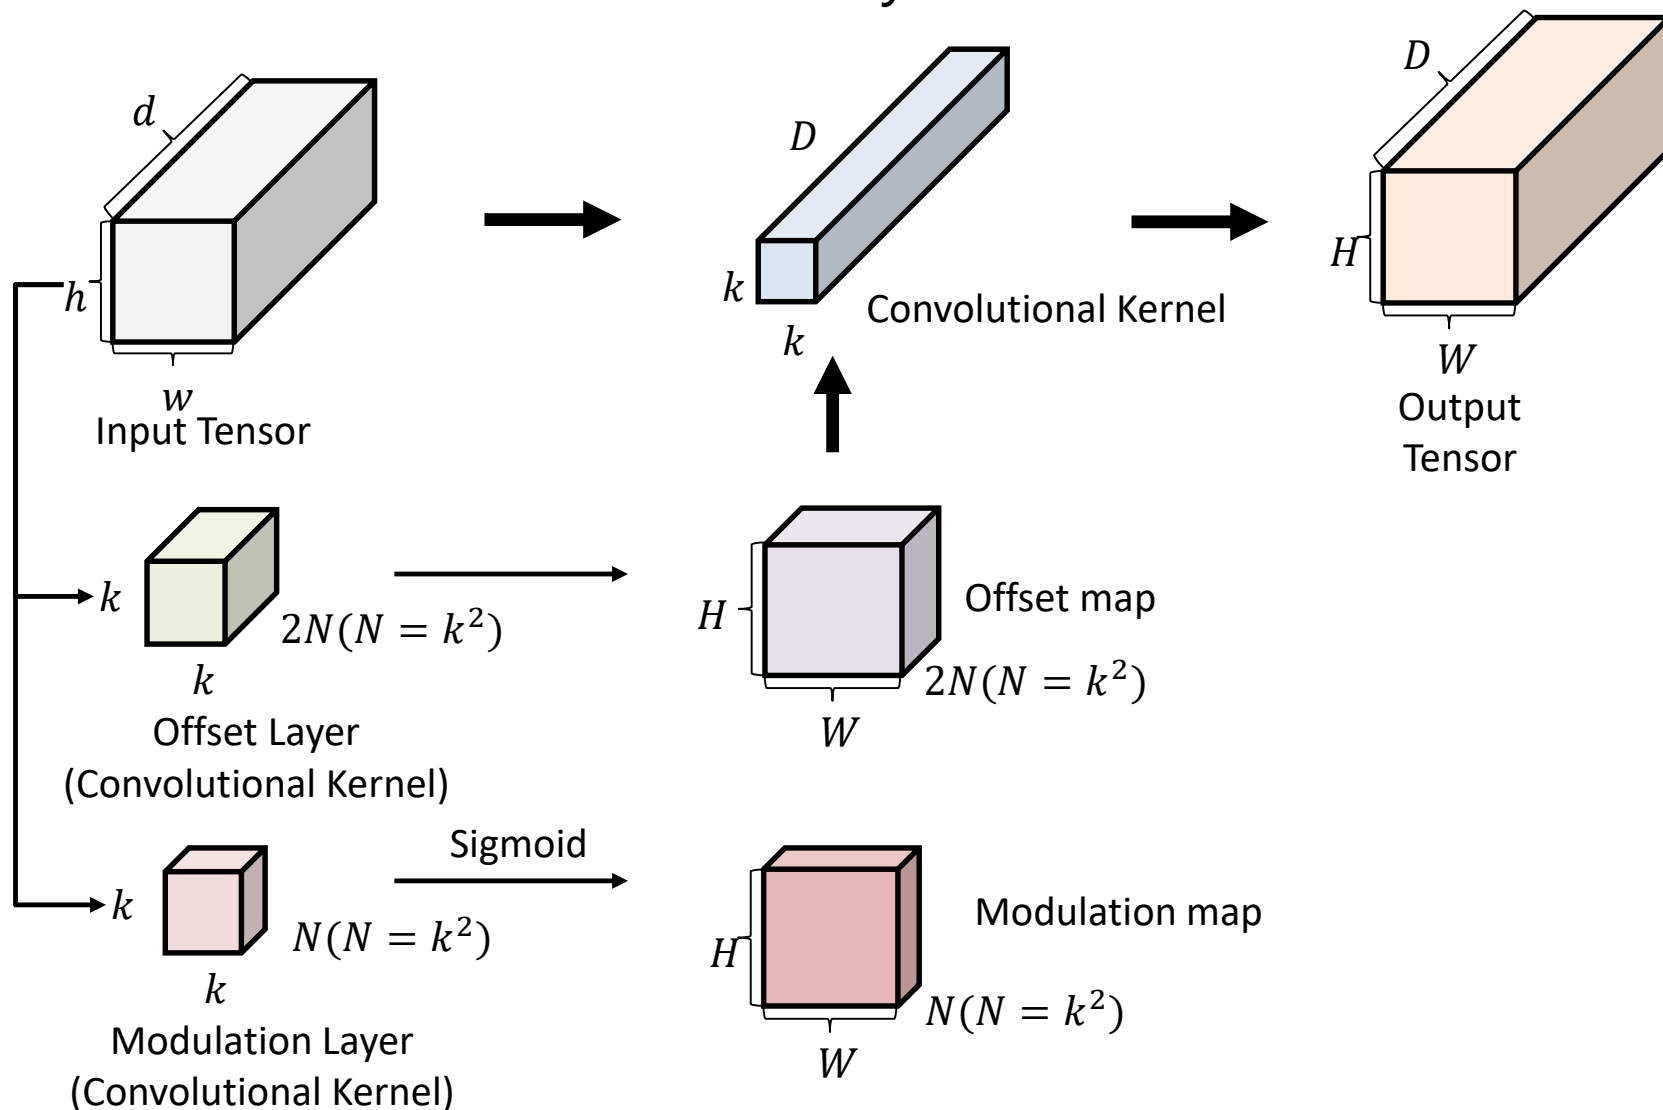
# Deformable Convolutional Layer v3

## ❖ Deformable Convolutional Layer v1



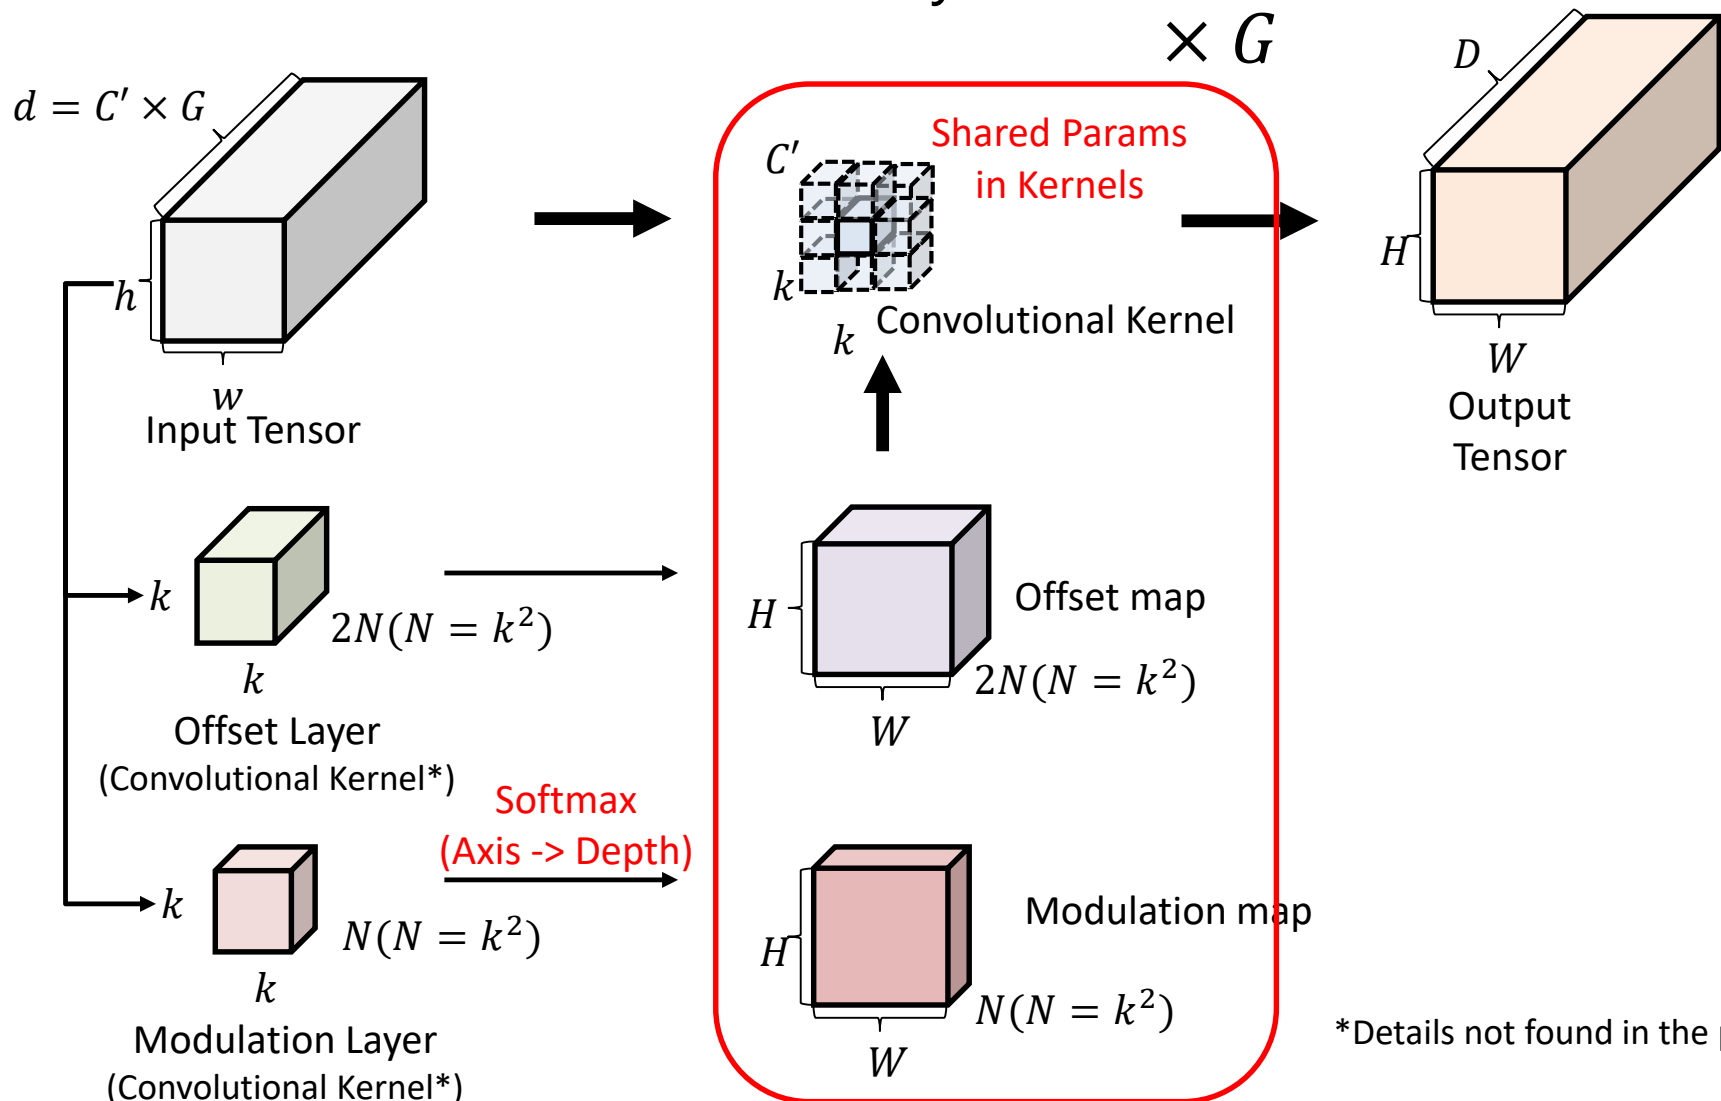
# Deformable Convolutional Layer v3

## ❖ Deformable Convolutional Layer v2

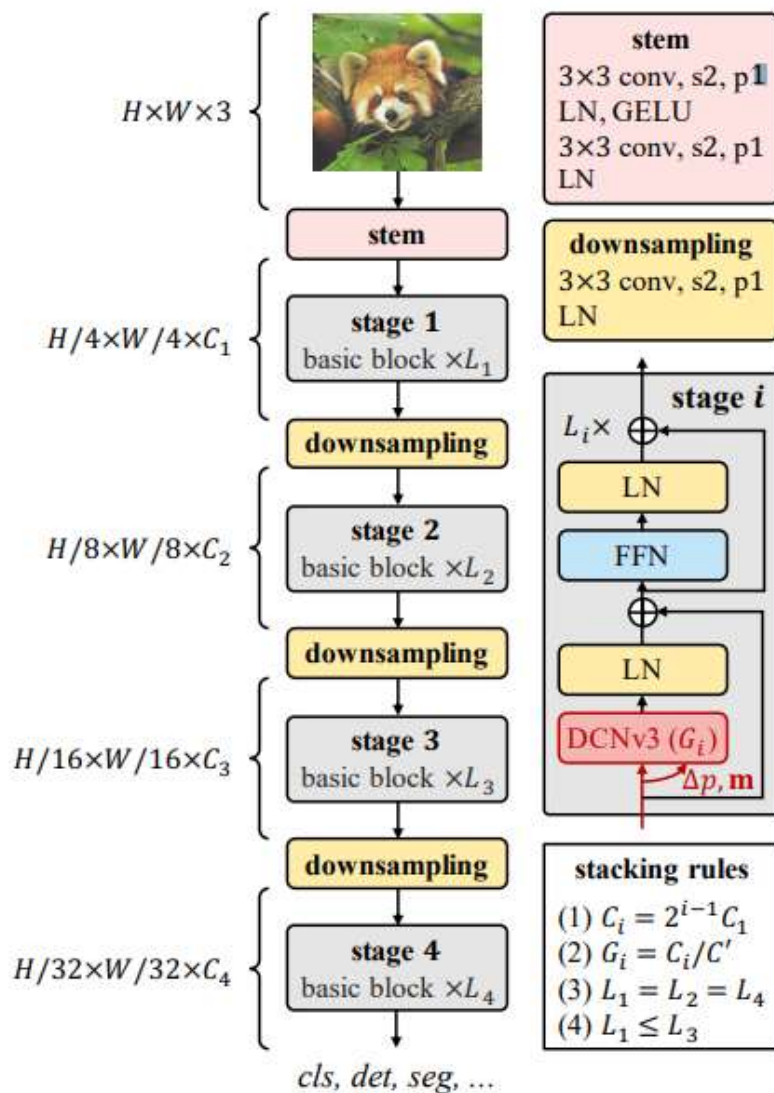


# Deformable Convolutional Layer v3

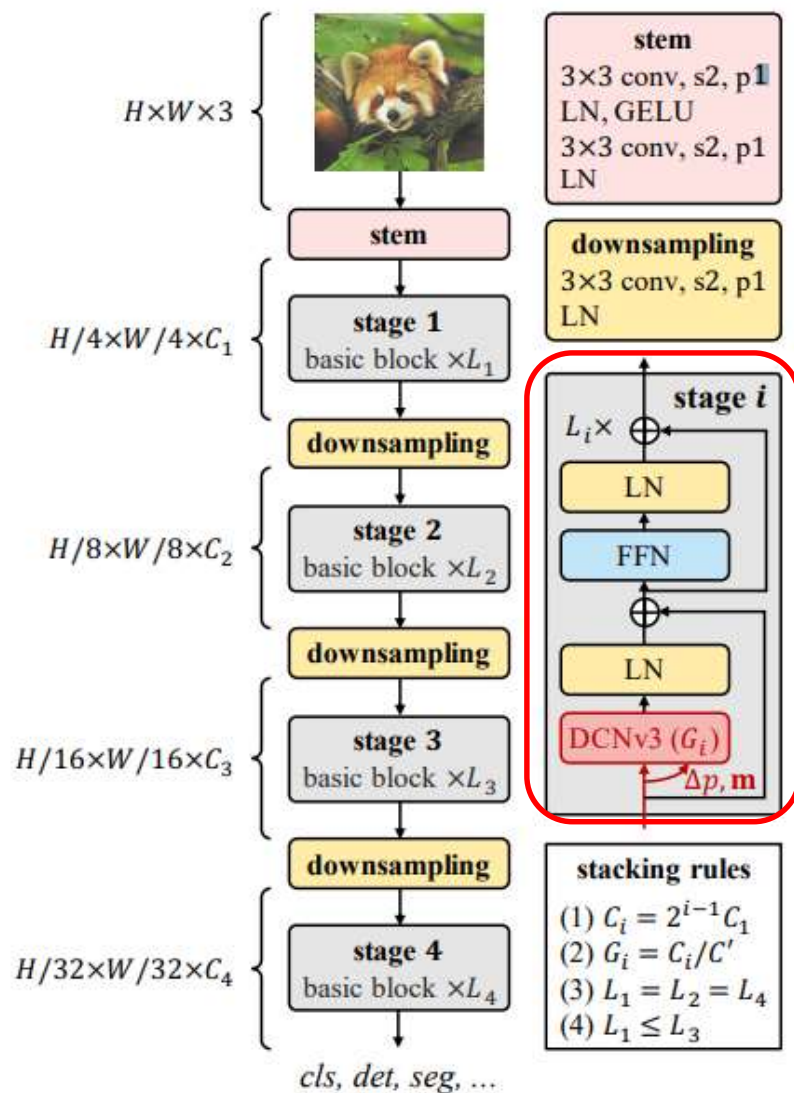
## ❖ Deformable Convolutional Layer v3



# InterImage Model

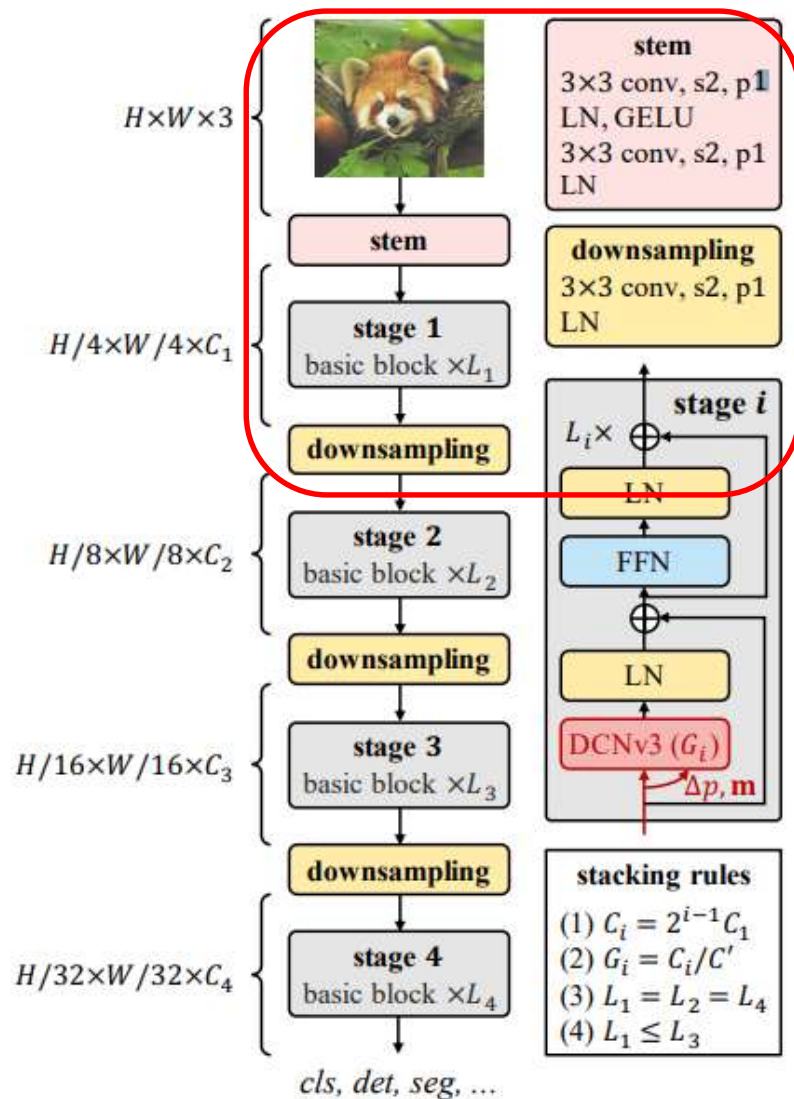


# InterImage Model



1. Basic Block

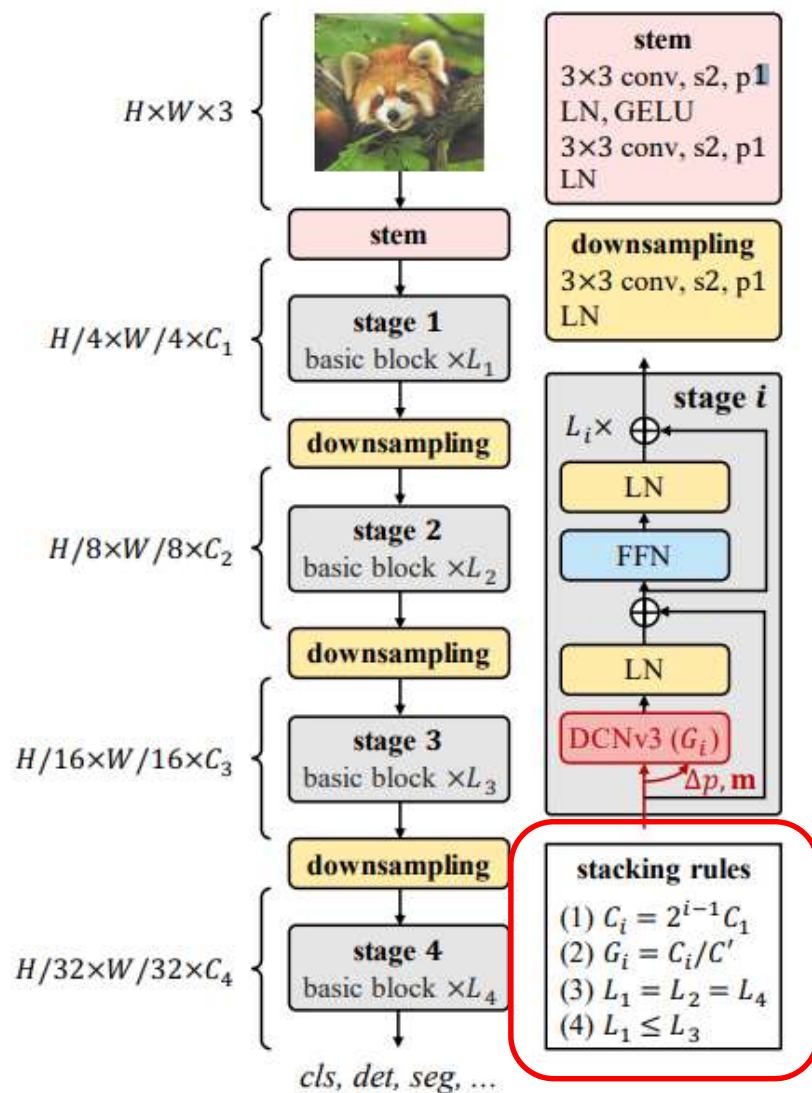
# InterImage Model



## 2. Stem layer & Downsampling



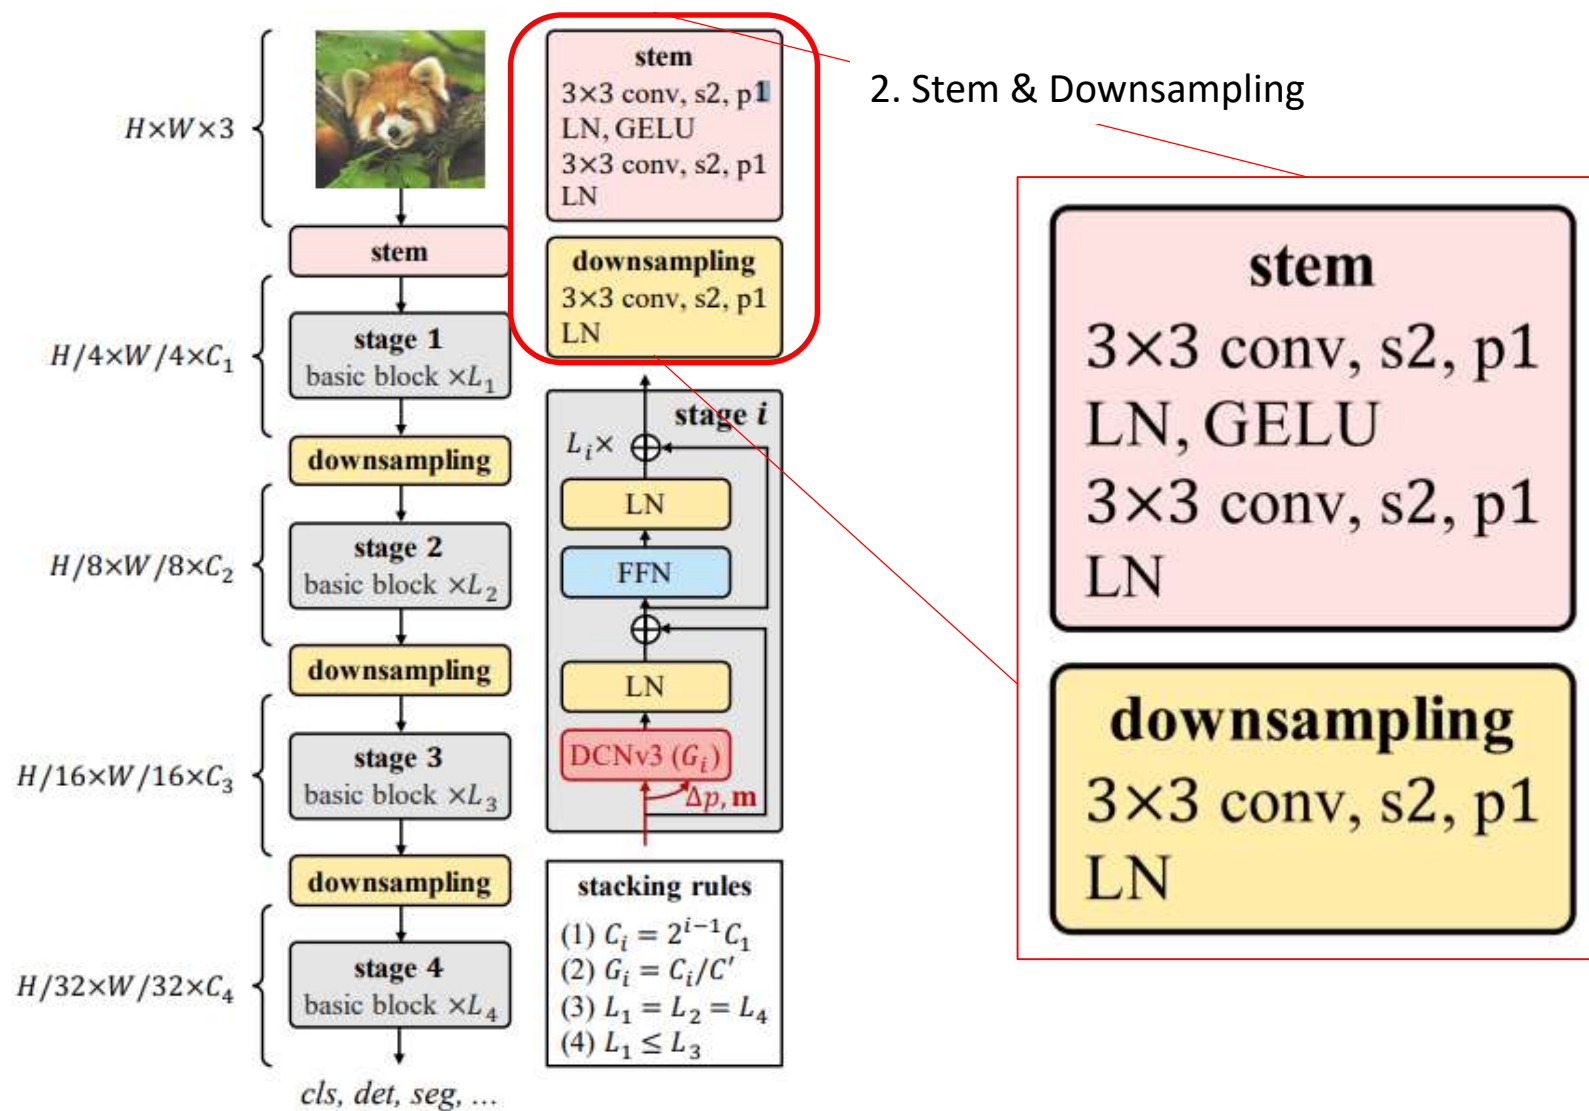
# InterImage Model



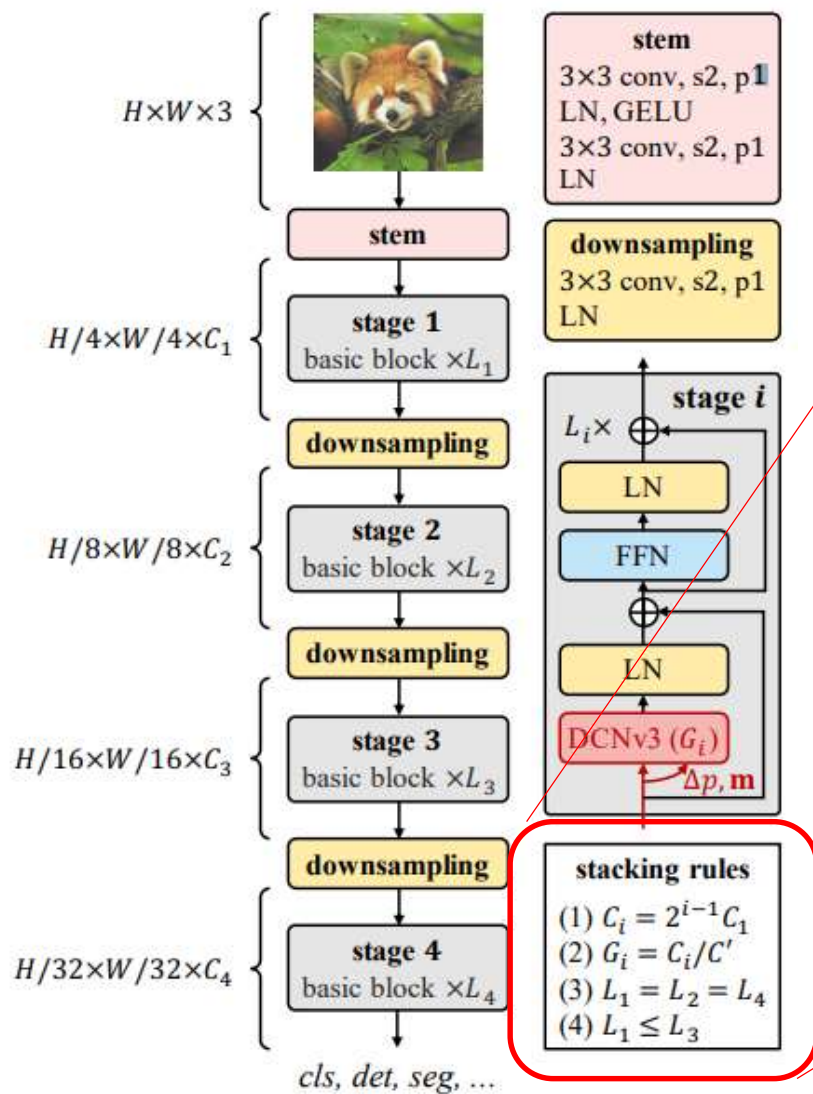
## 3. Stacking Rules



# InterImage Model



# InterImage Model



## 3. Stacking Rules

### stacking rules

- (1)  $C_i = 2^{i-1} C_1$
- (2)  $G_i = C_i / C'$
- (3)  $L_1 = L_2 = L_4$
- (4)  $L_1 \leq L_3$

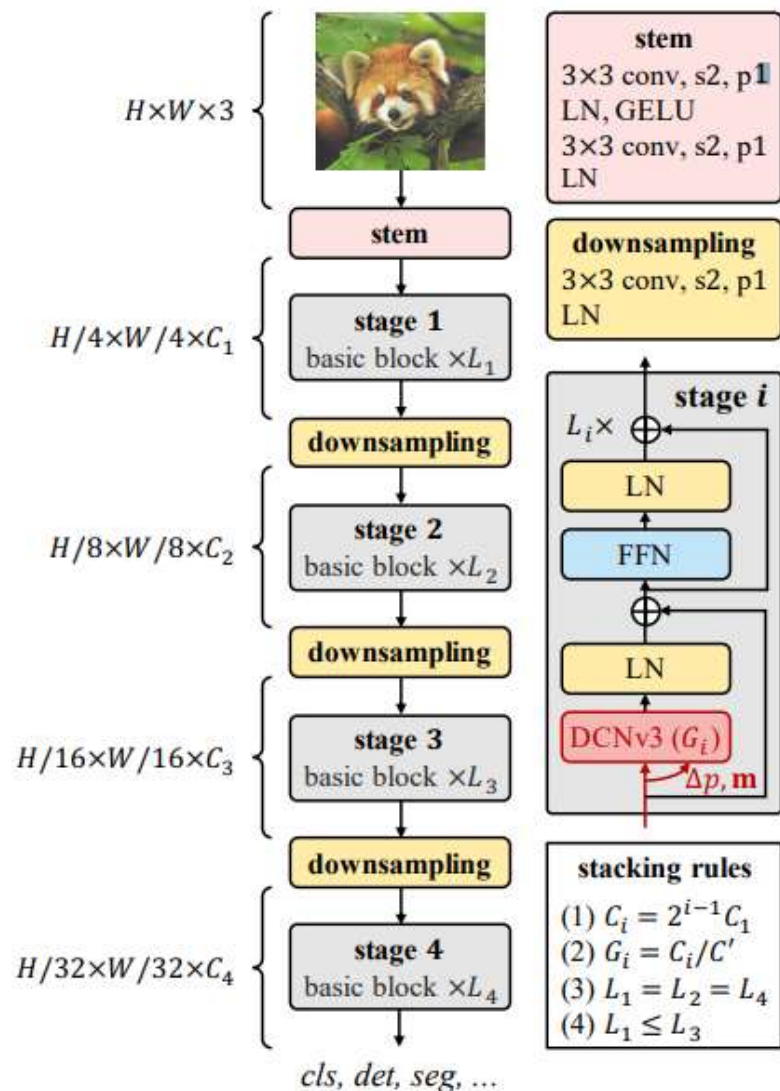
$C_i$ : the channel number of the  $i$ -th stage;

$G_i$ : the group number of the DCNv3 in the  $i$ -th stage;

$L_i$ : the number of basic blocks in the  $i$ -th stage.



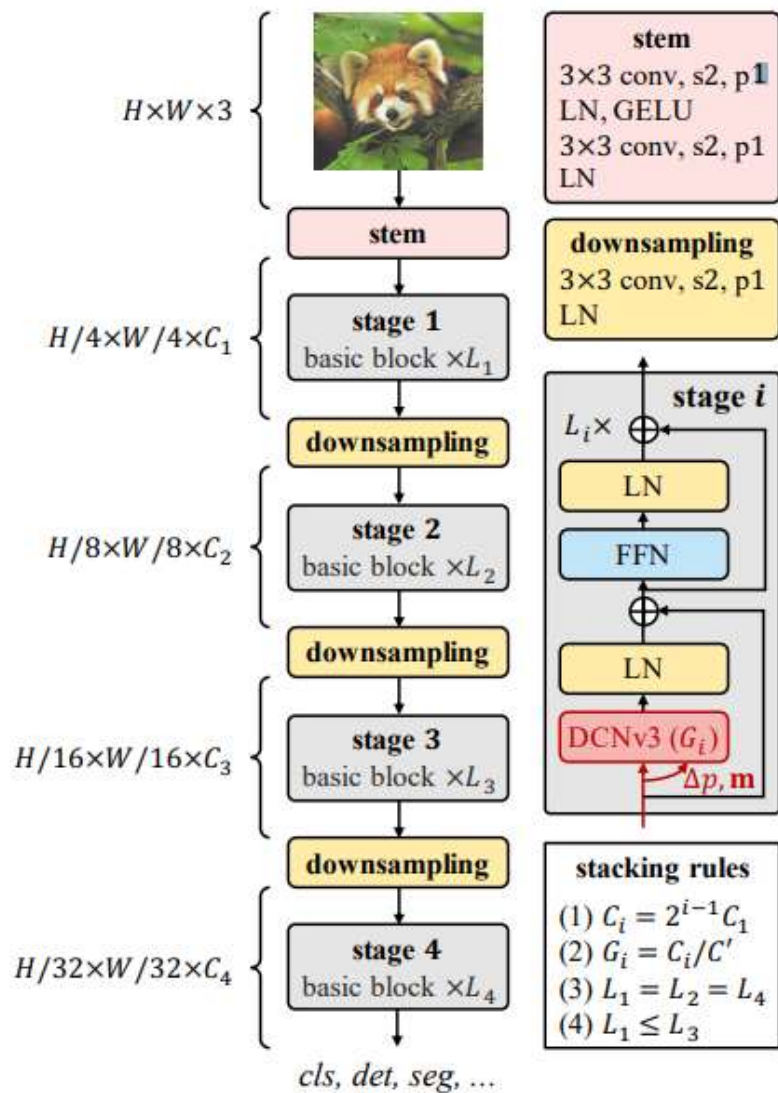
# InterImage Model



## 4. Scaling Rules

**Scaling rules.** Based on the optimal origin model under the aforementioned constraints, we further explore the parameter scaling rules inspired by [38]. Specifically, we consider two scaling dimensions: depth  $D$  (i.e.,  $3L_1 + L_3$ ) and width  $C_1$ , and scale the two dimensions using  $\alpha$ ,  $\beta$  and a composite factor  $\phi$ . The scaling rules can be written as:  $D' = \alpha^\phi D$  and  $C'_1 = \beta^\phi C_1$ , where  $\alpha \geq 1$ ,  $\beta \geq 1$ , and  $\alpha\beta^{1.99} \approx 2$ . Here, 1.99 is specific for InternImage and calculated by doubling the model width and keeping the depth constant. We experimentally find out that the best scaling setting is  $\alpha = 1.09$  and  $\beta = 1.36$ , and then we base on it to construct InternImage variants with different parameter scales, namely InternImage-T/S/B/L/XL, whose complexity is similar to those of ConvNeXt [21]. To further test the capability, we built a larger InternImage-H with 1 billion

# InterImage Model



## 5. Hyper-parameters for models of different scales

model name	$C_1$	$C'$	$L_{1,2,3,4}$	#params
InternImage-T (origin)	64	16	4, 4, 18, 4	30M
InternImage-S	80	16	4, 4, 21, 4	50M
InternImage-B	112	16	4, 4, 21, 4	97M
InternImage-L	160	16	5, 5, 22, 5	223M
InternImage-XL	192	16	5, 5, 24, 5	335M
InternImage-H	320	16	6, 6, 32, 6	1.08B

Table 1. **Hyper-parameters for models of different scales.** InternImage-T is the origin model, and -S/B/L/XL/H are scaled up from -T. “#params” denotes the number of parameters.



4 Stages Models are prevalent since DETR

- Swin Transformer
- MetaFormer



## Q & A

---

Q & A



# Experiments

## ❖ Image Classification (Tiny Model)

method	type	scale	#params	#FLOPs	acc (%)
DeiT-S [58]	T	$224^2$	22G	5G	79.9
PVT-S [10]	T	$224^2$	25M	4G	79.8
Swin-T [2]	T	$224^2$	29M	5G	81.3
CoAtNet-0 [20]	T	$224^2$	25M	4G	81.6
CSwin-T [12]	T	$224^2$	23M	4G	82.7
PVTv2-B2 [11]	T	$224^2$	25M	4G	82.0
DeiT III-S [65]	T	$224^2$	22M	5G	81.4
SwinV2-T/8 [16]	T	$256^2$	28M	6G	81.8
Focal-T [66]	T	$224^2$	29M	5G	82.2
ConvNeXt-T [21]	C	$224^2$	29M	5G	82.1
SLaK-T [29]	C	$224^2$	30M	5G	82.5
HorNet-T [44]	C	$224^2$	23M	4G	83.0
InternImage-T (ours)	C	$224^2$	30M	5G	83.5



# Experiments

## ❖ Image Classification (Large Model)

method	type	scale	#params	#FLOPs	acc (%)
ViT-G/14 <sup>#</sup> [30]	T	518 <sup>2</sup>	1.84B	5160G	90.5
CoAtNet-6 <sup>#</sup> [20]	T	512 <sup>2</sup>	1.47B	1521G	90.5
CoAtNet-7 <sup>#</sup> [20]	T	512 <sup>2</sup>	2.44B	2586G	90.9
Florence-CoSwin-H <sup>#</sup> [59]	T	—	893M	—	90.0
SwinV2-G <sup>#</sup> [16]	T	640 <sup>2</sup>	3.00B	—	90.2
RepLKNet-XL <sup>#</sup> [22]	C	384 <sup>2</sup>	335M	129G	87.8
BiT-L-ResNet152x4 <sup>#</sup> [64]	C	480 <sup>2</sup>	928M	—	87.5
InternImage-H <sup>#</sup> (ours)	C	224 <sup>2</sup>	1.08B	188G	88.5
InternImage-H <sup>#</sup> (ours)	C	640 <sup>2</sup>	1.08B	1478G	89.2

Table 2. **Image classification performance on the ImageNet validation set.** “type” refers to model type, where “T” and “C” denote transformer and CNN, respectively. “scale” is the input scale. “acc” is the top-1 accuracy. “<sup>‡</sup>” indicates the model is pre-trained on ImageNet-22K [31]. “<sup>#</sup>” indicates pretraining on extra large-scale private dataset such as JFT-300M [67], FLD-900M [59], or the joint public dataset in this work.

# Experiments

## ❖ Object Detection & Instance Segmentation

method	#params	#FLOPs	Mask R-CNN 1× schedule						Mask R-CNN 3×+MS schedule					
			AP <sup>b</sup>	AP <sup>b</sup> <sub>50</sub>	AP <sup>b</sup> <sub>75</sub>	AP <sup>m</sup>	AP <sup>m</sup> <sub>50</sub>	AP <sup>m</sup> <sub>75</sub>	AP <sup>b</sup>	AP <sup>b</sup> <sub>50</sub>	AP <sup>b</sup> <sub>75</sub>	AP <sup>m</sup>	AP <sup>m</sup> <sub>50</sub>	AP <sup>m</sup> <sub>75</sub>
Swin-T [2]	48M	267G	42.7	65.2	46.8	39.3	62.2	42.2	46.0	68.1	50.3	41.6	65.1	44.9
ConvNeXt-T [21]	48M	262G	44.2	66.6	48.3	40.1	63.3	42.8	46.2	67.9	50.8	41.7	65.0	44.9
PVTv2-B2 [11]	45M	309G	45.3	67.1	49.6	41.2	64.2	44.4	47.8	69.7	52.6	43.1	66.8	46.7
ViT-S [9,68]	48M	353G	44.7	65.8	48.3	39.9	62.5	42.8	48.2	69.7	52.5	42.8	66.4	45.9
InternImage-T (ours)	49M	270G	47.2	69.0	52.1	42.5	66.1	45.8	49.1	70.3	54.0	43.7	67.3	47.1
Swin-S [2]	69M	354G	44.8	66.6	48.9	40.9	63.4	44.2	48.2	69.8	52.8	43.2	67.0	46.1
ConvNeXt-S [21]	70M	348G	45.4	67.9	50.0	41.8	65.2	45.1	47.9	70.0	52.7	42.9	66.9	46.2
PVTv2-B3 [11]	65M	397G	47.0	68.1	51.7	42.5	65.7	45.7	48.4	69.8	53.3	43.2	66.9	46.7
InternImage-S (ours)	69M	340G	47.8	69.9	52.8	43.3	67.1	46.7	49.7	71.1	54.5	44.4	68.5	47.8
Swin-B [2]	107M	496G	46.9	—	—	42.3	—	—	48.6	70.0	53.4	43.3	67.1	46.7
ConvNeXt-B [21]	108M	486G	47.0	69.4	51.7	42.7	66.3	46.0	48.5	70.1	53.3	43.5	67.1	46.7
PVTv2-B5 [11]	102M	557G	47.4	68.6	51.9	42.5	65.7	46.0	48.4	69.2	52.9	42.9	66.6	46.2
ViT-B [9,68]	120M	781G	47.0	68.2	51.4	41.8	65.1	44.9	49.6	70.6	54.0	43.6	67.7	46.9
InternImage-B (ours)	115M	501G	48.8	71.0	53.9	44.0	67.8	47.5	50.3	71.4	55.3	44.8	68.7	48.0

method	#param	#FLOPs	Cascade Mask R-CNN 1× schedule						Cascade Mask R-CNN 3×+MS schedule					
			AP <sup>b</sup>	AP <sup>b</sup> <sub>50</sub>	AP <sup>b</sup> <sub>75</sub>	AP <sup>m</sup>	AP <sup>m</sup> <sub>50</sub>	AP <sup>m</sup> <sub>75</sub>	AP <sup>b</sup>	AP <sup>b</sup> <sub>50</sub>	AP <sup>b</sup> <sub>75</sub>	AP <sup>m</sup>	AP <sup>m</sup> <sub>50</sub>	AP <sup>m</sup> <sub>75</sub>
Swin-L <sup>‡</sup> [2]	253M	1382G	51.8	71.0	56.2	44.9	68.4	48.9	53.9	72.4	58.8	46.7	70.1	50.8
ConvNeXt-L <sup>‡</sup> [21]	255M	1354G	53.5	72.8	58.3	46.4	70.2	50.2	54.8	73.8	59.8	47.6	71.3	51.7
RepLKNet-31L <sup>‡</sup> [22]	229M	1321G	—	—	—	—	—	—	53.9	72.5	58.6	46.5	70.0	50.6
HorNet-L <sup>‡</sup> [44]	259M	1358G	—	—	—	—	—	—	56.0	—	—	48.6	—	—
InternImage-L <sup>‡</sup> (ours)	277M	1399G	54.9	73.8	59.6	47.7	71.3	52.4	56.0	74.7	61.3	48.4	72.2	53.0
ConvNeXt-XL <sup>‡</sup> [21]	407M	1898G	53.6	72.9	58.5	46.5	70.3	50.5	55.2	74.2	59.9	47.7	71.6	52.2
InternImage-XL <sup>‡</sup> (ours)	387M	1782G	55.3	74.5	60.2	48.0	72.0	52.4	56.2	74.9	61.7	48.8	72.6	53.8

Table 3. **Object detection and instance segmentation performance on COCO val2017.** The FLOPs are measured with 1280×800 inputs. AP<sup>b</sup> and AP<sup>m</sup> represent box AP and mask AP, respectively. “MS” means multi-scale training.



# Experiments

## ❖ Object Detection & Instance Segmentation

method	detector	#params	AP <sup>b</sup>	
			val2017	test-dev
Swin-L <sup>†</sup> [2]	HTC++ [2]	284M	58.0	58.7
Swin-L [2]	DyHead [72]	213M	56.2	58.4
ViT-L <sup>†</sup> [9]	ViT-Adapter [68]	401M	60.5	60.9
Swin-L <sup>†</sup> [2]	Soft-Teacher [73]	284M	60.7	61.3
Swin-L <sup>†</sup> [2]	DINO [74]	218M	63.2	63.3
FocalNet-H <sup>†</sup> [75]	DINO [74]	746M	64.2	64.3
ViT-Huge [76]	Group-DETRv2 [76]	629M	—	64.5
Florence-CoSwin-H <sup>#</sup> [59]	DyHead [72]	637M	62.0	62.4
SwinV2-G <sup>#</sup> [16]	HTC++ [2]	3.00B	62.5	63.1
BEiT-3 <sup>#</sup> [17]	ViTDet [77]	1.90B	—	63.7
FD-SwinV2-G <sup>#</sup> [26]	HTC++ [2]	3.00B	—	64.2
InternImage-XL <sup>†</sup> (ours)	DINO [74]	602M	64.2	64.3
InternImage-H <sup>#</sup> (ours)	DINO [74]	2.18B	65.0	65.4

Table 4. Comparison of the state-of-the-art detectors on COCO val2017 and test-dev.

# Experiments

## ❖ Semantic Segmentation

method	crop size	#params	#FLOPs	mIoU (SS)	mIoU (MS)
Swin-T [2]	512 <sup>2</sup>	60M	945G	44.5	45.8
ConvNeXt-T [21]	512 <sup>2</sup>	60M	939G	46.0	46.7
SLaK-T [29]	512 <sup>2</sup>	65M	936G	47.6	—
InternImage-T (ours)	512 <sup>2</sup>	59M	944G	47.9	48.1
Swin-S [2]	512 <sup>2</sup>	81M	1038G	47.6	49.5
ConvNeXt-S [21]	512 <sup>2</sup>	82M	1027G	48.7	49.6
SLaK-S [29]	512 <sup>2</sup>	91M	1028G	49.4	—
InternImage-S (ours)	512 <sup>2</sup>	80M	1017G	50.1	50.9
Swin-B [2]	512 <sup>2</sup>	121M	1188G	48.1	49.7
ConvNeXt-B [21]	512 <sup>2</sup>	122M	1170G	49.1	49.9
RepLKNet-31B [22]	512 <sup>2</sup>	112M	1170G	49.9	50.6
SLaK-B [29]	512 <sup>2</sup>	135M	1172G	50.2	—
InternImage-B (ours)	512 <sup>2</sup>	128M	1185G	50.8	51.3
Swin-L <sup>†</sup> [2]	640 <sup>2</sup>	234M	2468G	52.1	53.5
RepLKNet-31L <sup>†</sup> [22]	640 <sup>2</sup>	207M	2404G	52.4	52.7
ConvNeXt-L <sup>†</sup> [21]	640 <sup>2</sup>	235M	2458G	53.2	53.7
ConvNeXt-XL <sup>†</sup> [21]	640 <sup>2</sup>	391M	3335G	53.6	54.0
InternImage-L <sup>†</sup> (ours)	640 <sup>2</sup>	256M	2526G	53.9	54.1
InternImage-XL <sup>†</sup> (ours)	640 <sup>2</sup>	368M	3142G	55.0	55.3
SwinV2-G <sup>#</sup> [16]	896 <sup>2</sup>	3.00B	—	—	59.9
InternImage-H <sup>#</sup> (ours)	896 <sup>2</sup>	1.12B	3566G	59.9	60.3
BEiT-3 <sup>#</sup> [17]	896 <sup>2</sup>	1.90B	—	—	62.8
FD-SwinV2-G <sup>#</sup> [26]	896 <sup>2</sup>	3000	—	—	61.3
InternImage-H <sup>#</sup> (ours) + Mask2Former [80]	896 <sup>2</sup>	1.31B	4635G	62.5	62.9

Table 5. **Semantic segmentation performance on the ADE20K validation set.** The FLOPs are measured with 512×2048, 640×2560, or 896×896 inputs according to the crop size. “SS” and “MS” means single-scale and multi-scale testing, respectively.



# 이미지처리팀 리뷰 의견

---

## ❖ Deformable Conv V3에 대한 분석이 부재

### ▣ Ablation Study 부재

- ResNet 등 기존 Conv 기반 모델에서 성능향상을 가지는지 확인해줬으면 정말 좋았을 듯

### ▣ Deformable Conv V3의 장점에 대한 정성적인 분석 부재

- Convolution을 Group으로 나누면서 생기는 단일 레이어의 다양한 Offset Map의 장점에 대한 시각화 자료가 있었으면 좋았을 듯

### ▣ Inductive Bias를 줄일 수 있었다는 주장에 대한 근거자료 부재

## ❖ 코드가 아직 공개되지 않아 정확한 검증은 어려움



## Q & A

---

Q & A